

AI Scan and IVR

What has been delivered, and what belongs to the next phase

This note covers the two AI subjects you raised. They sit in very different places, and separating them is the whole purpose of the document.

The AI Scan was in the agreed three milestone plan. It is built, tested and deployed. Every part of it that was our responsibility to deliver is finished, and the only thing standing between it and live reports is a provider decision and an API key, both of which sit with you. That position is not new: the Milestone 2 tracker recorded it in those terms at the time.

The AI Receptionist, the Scheduling Agent, AI Sales Outbound and IVR were never in the plan. What exists today is a design preview, and it is labelled as sample data inside the product so that nobody mistakes it for a working system. These are the natural content of a next phase, and the second half of this note sets out what building them would involve.

35

Questions across
five dimensions

3

Scan tiers
live

46

Automated tests
passing

1

Item outstanding,
and it is a key

Part one: the AI Scan

The AI Scan is the public facing assessment that turns a prospect's answers into a scored operational report. It runs end to end today. What follows is each piece of it, what that piece does, and how we checked it rather than assumed it.

Intake and questionnaire

A thirty five question assessment across five dimensions

DELIVERED

Operational maturity, automation readiness, infrastructure maturity, scalability readiness and AI opportunity potential. The questions are held in one shared definition used by both the website and the scoring engine, so the two can never disagree about what was asked.

Three tiers, each with its own depth

DELIVERED

The free tier asks seven questions, growth asks eighteen and elite asks all thirty five. A visitor is served the correct set automatically, and the server independently refuses a submission that is missing an answer the chosen tier requires. It also refuses any answer carrying a question identifier that is not in the bank, so the record cannot be polluted by a crafted request.

Lead capture wired into the CRM

DELIVERED

Every completed questionnaire creates a lead record alongside the scan itself, so a prospect appears in the pipeline whether or not their report finishes. Campaign source and campaign name are carried through where present.

Anti abuse measures on a public endpoint

DELIVERED

The scan is open to the internet, so it is rate limited per network address and carries a hidden field that automated form fillers complete and humans do not. A submission that trips it is accepted silently rather than rejected, which tells the sender nothing useful.

Scoring

A deterministic scoring layer, independent of any AI provider

DELIVERED

Every answer option carries a numeric score, and a raw score per dimension is computed in our own code before the model is involved at all. Those raw scores are given to the model as part of its input, and they are also what the report falls back to for any dimension the model does not return a usable number for. The result is that a report is still produced, and still grounded in the actual answers, even when the model output is incomplete.

An AI enrichment layer on top of it

DELIVERED

The model is asked to produce written reasoning, detected opportunities and a roadmap, constrained to a fixed response schema so the output can be relied on rather than parsed hopefully. Where it returns a score for a dimension, that score is used, clamped to a valid range. Where it does not, the deterministic score computed beforehand stands in. Written reasoning is treated the same way, with a plain fallback sentence rather than an empty section.

Provider neutrality, deliberately

DELIVERED

The integration speaks the OpenAI chat completions format, which OpenAI, Azure OpenAI, Google Gemini and several others all accept directly. Switching provider is a configuration change, not a rebuild. There is a request timeout and explicit handling of rate limit responses, neither of which existed before this work.

The report

A structured report, generated as a PDF

DELIVERED

Five sections: executive summary, dimension breakdown, detected opportunities, operational roadmap and important notes. The PDF is produced in our own code without a headless browser, which is what makes it safe to run on a normal deployment.

Private storage with token gated access

DELIVERED

Reports are stored privately rather than at a guessable address. A recipient reaches their own report through a random token issued to them, and the link they receive is short lived and signed. One person's token does not reach another person's report.

Administration

Retrigger a failed or incomplete scan

DELIVERED

An administrator can rerun the scoring on an existing scan rather than asking the prospect to fill the form again. If the engine fails, the failure is recorded against the scan with its reason rather than being swallowed.

Promote a scan into a client report

DELIVERED

A completed scan can be converted into a formal report attached to a client organisation. The scan itself is anonymous public intake, so this is the deliberate bridge between a prospect's assessment and a named client's record.

Honest configuration reporting in the admin console

DELIVERED

The console shows whether an AI provider is configured, derived from the actual environment rather than from a value someone typed in. It also carries the retention position for scan data, recorded truthfully as having no automated deletion configured rather than implying an enforcement that does not exist.

How this was checked

Forty six automated tests cover this subsystem specifically, and all of them pass. They exercise the question bank and tier selection, the deterministic scoring, the AI enrichment path with a stubbed provider, the PDF generation, the public endpoints including the rate limit and the hidden field, the administrator retrigger and the bridge into client reports.

Beyond the test suite, we ran a real scan against the live deployment rather than relying on tests alone. The lead record was created, the scan record was created, the questionnaire was validated and stored, and the process reached the scoring step. It stopped there, for the reason on the next page.

The one item still open, and it is not code

The AI Scan needs a model provider to be chosen and an API key to be installed. Until that happens the scoring step has nothing to call, so reports do not complete.

This is not a gap that appeared recently. It was recorded in the Milestone 2 tracker at the time the work was done, in these words:

MILESTONE 2, SECTION 2.2

Production activation of the LLM provider is BLOCKED (client decisions + credentials). Everything code-side is done; this is a client decision + credential, not a code gap.

Two things are needed from your side, and the second takes a minute once the first is settled.

- A decision on which provider to use, and on what that provider is contractually permitted to do with the data that passes through it. That second half matters and should be settled before a key is installed rather than afterwards, because it is far harder to unwind later.
- The API key itself. We have already changed the code so that a key placed in the environment's existing OpenAI slot is enough on its own to activate the scan. Previously the code read a differently named setting, which meant a key could sit there populated and unread. That is fixed.

Nothing else is required from us. The moment a key is in place, scans complete, reports generate, PDFs are produced and stored, and the administrator tools work against real results.

A note on what we are not claiming

We are not describing the AI Scan as proven in production, because a report has not yet been generated against a live provider. Everything up to that call is exercised and working, and the call itself is covered by tests against a stubbed provider. We would rather set that out plainly than describe it as finished and have you discover the distinction yourself.

Part two: voice, IVR and the agents

The AI Receptionist, the Scheduling Agent, AI Sales Outbound and IVR are a different matter. None of them appear in Milestone 1, 2 or 3. Every task in that plan carries a reference number, and none of these has one, because no workstream ever contained them.

Where the agreed documents mention the subject at all, it is to record that it was being left out rather than built:

MILESTONE 2 TRACKER

no real backing system exists for any of these: alerting, workflow automation, AI voice/IVR, infra health-check, campaign management

MILESTONE 2, SECTION 2.4

Campaigns/Agents/Automations/Analytics-funnel deliberately deferred: genuinely new subsystems, out of bounded scope.

What exists today

There is a screen in the administrator console titled AI Agents and IVR. It shows what these products would look like, carrying placeholder figures. It is marked with a sample data badge in the interface, and during Milestone 2 we removed the misleading live indicators that had been sitting on it, precisely so that it could not be mistaken for a working system. We mention this because it is the thing most likely to create a misunderstanding about what is built.

One further clarification worth making, because it is easy to assume otherwise. The platform does hold Twilio credentials, and they do work. They are connected to text messaging only, and are used to deliver multi factor authentication codes. Nothing in the system has ever been connected to voice.

What a next phase would need to build

A telephony foundation

NEXT PHASE

Numbers, inbound and outbound call handling, call state, and recording storage with its own retention position. All four products sit on this one foundation, which is the main reason we would recommend scoping them together rather than one at a time. Quoting them separately means paying for this groundwork more than once.

Speech handling

NEXT PHASE

Speech to text on the way in and text to speech on the way out, with a decision on provider and on language coverage. The platform already supports ten written languages, and whether the voice products need to match that is a real cost decision rather than a detail.

AI Receptionist

NEXT PHASE

The conversational layer that answers, understands intent, and knows when it is out of its depth. Includes the escalation path to a person, which is the part that decides whether callers trust it.

Scheduling Agent

NEXT PHASE

Booking and rescheduling in conversation. The booking engine underneath is already real and working, so this is the conversational layer rather than the calendar. That is the one place where the existing platform meaningfully reduces the work.

IVR

NEXT PHASE

Menu driven routing, with the route definitions editable by an administrator rather than fixed in code.

AI Sales Outbound

NEXT PHASE

Campaign definition, sequencing, targeting and conversion tracking, with the compliance constraints that outbound calling carries in each market you operate in.

Recording and transcription, independently controllable

NEXT PHASE

You asked that these two remain separately switchable rather than technically inseparable. We agree, and we have recorded it as a design constraint. It is worth stating explicitly in the phase scope, because it is significantly cheaper to build in from the start than to retrofit.

Summary

The AI Scan

In the agreed plan, and delivered. Intake, questionnaire, deterministic scoring, AI enrichment, report generation, PDF output, private storage, token gated access, administrator retrigger and the bridge into client reports are all built and covered by forty six passing tests. One item remains, it is a provider decision and a key, and it sits with you. Your own Milestone 2 tracker recorded that position at the time.

Voice, IVR and the agents

Not in the agreed plan, and not started. What exists is a design preview, labelled as sample data inside the product. These four products share one telephony and speech foundation that does not exist yet, and we would recommend scoping them as a single phase so that foundation is costed once rather than four times.

What would help most

- Send the AI provider key, and tell us your position on what that provider may do with the data.
- Confirm whether the voice products should be scoped as a phase of their own. If so, we will put a proper proposal together covering all four.
- Tell us which capability the companies you mentioned are actually waiting on. That single answer changes the order we work in more than any other piece of information you could give us.

Scope basis: IO_SKY_Three_Milestone_Plan.pdf, tracked through PHASE1_CHECKLIST.md, MILESTONE2_PROGRESS.md and MILESTONE3_CHECKLIST.md. Quotations are taken from those documents, with punctuation normalised. Test counts and behaviour described here were re-run and re-verified on 14 September 2026.