

AI Scan and IVR

What the three agreed milestones covered, and what the new specifications require

This note answers one question directly: if the AI Scan and the IVR already exist in the platform, why do the new specifications form a separate phase of work rather than completion of the work already paid for.

It is a fair question and it deserves a precise answer rather than an assertion. The answer is different for each of the two, so they are handled separately below. Both answers are drawn from the specifications IO SKY has supplied and from the project tracking documents written at the time, not from our recollection.

In short: the IVR does not exist, and the AI Scan that exists is a different product from the AI Scan the new specifications define. The new specifications say so themselves, in their own acceptance tests.

8

Specification
documents received

133

Pages of
requirements

302

Implementation
tasks derived

0

Of them in the
agreed plan

Part one: the IVR

This one is quick, because there is nothing to compare. No telephony exists in the platform. No call has ever been received or placed. There is no voice integration of any kind.

What exists is a screen in the administrator console titled AI Agents and IVR. It displays example agents with example call volumes. It is marked with a sample data badge in the interface, and that badge was added deliberately during Milestone 2 precisely so that it could not be mistaken for a working system. At the same time we removed the misleading live indicators that had been sitting on it.

The Milestone 2 tracker recorded this at the time, in a list of screens that were given the sample data badge because no real system stood behind them:

MILESTONE 2 TRACKER, WRITTEN AUGUST 2026

Disclosed with the sample-data badge (no real backing system exists for any of these: alerting, workflow automation, AI voice/IVR, infra health-check, campaign management ...)

The same tracker recorded the commercial position on this family of features in the section that defines what Milestone 2 delivered:

MILESTONE 2, SECTION 2.4 STATUS LINE

Campaigns/Agents/Automations/Analytics-funnel deliberately deferred: genuinely new subsystems, out of bounded scope.

One point worth clearing up, because it is the most likely source of the impression that telephony exists. The platform does hold Twilio credentials and they do work. They are connected to text messaging only, and they are used to deliver multi-factor authentication codes when someone signs in. Nothing has ever been connected to voice.

What the IVR specification now requires

The IVR and Telephony Functional Specification runs to 47 sections and carries a 37-item acceptance test matrix. Among its mandatory requirements: a real DTMF fallback menu in Dutch and English, a queue with configurable limits and abandon tracking, availability-aware routing with primary, secondary and fallback destinations, a compact handoff summary delivered to the operator before they answer, provider-confirmed transfer outcomes across nine distinct states, callback-first fallback with caller ID awareness, voicemail, recording as an operational telephony state rather than a database flag, and ten genuinely simultaneous calls proven in an acceptance test that the specification states is explicitly not a capacity ceiling.

None of this is an extension of a sample-data screen. It is a telephony platform.

Part two: the AI Scan

This one is a genuinely reasonable question, because an AI Scan does exist and it does work. So the answer has to be more careful than pointing at a sample data badge.

The short version: two different products share the same name. The one that exists was built as a lead-generation assessment. The one the new specifications define is a paid professional consulting product with human expert review, controlled publication and an audiovisual executive briefing. The new specification is explicit that the first cannot become the second.

The specification grades the existing implementation as a failure, twice

This is not our interpretation. The Master Specification contains 30 numbered acceptance tests. Two of them are written as conditions that must FAIL, and the AI Scan running in production today meets both of them exactly.

MASTER SPECIFICATION, TST-001 FORM TEST

If experience is effectively customer fills long form, then AI generates PDF: FAIL

MASTER SPECIFICATION, TST-004 FIXED SEQUENCE TEST

Hidden fixed 35-question sequence with cosmetic variation: FAIL

The AI Scan in production is a fixed questionnaire that produces a PDF. The file `shared/aiScanQuestionnaire.ts` contains exactly 35 questions. Both tests are graded FAIL rather than marked as needing improvement, which means the existing engine cannot be adjusted into compliance. The specification reinforces this in its opening product definition:

MASTER SPECIFICATION, PRD-001

The AI Scan is a premium adaptive professional assessment and decision-support product. It MUST NOT function or feel like a static questionnaire, survey, lead-generation quiz, fixed question tree, maturity calculator, generic chatbot, long form or questionnaire followed by an automatically generated PDF.

Four further contradictions, verified against the live database

These are not judgement calls. Each was checked in the running code rather than assumed.

The commercial model

IN PRODUCTION TODAY

Three tiers named free, growth and elite. The free tier is genuinely free and asks seven questions. No payment processor is connected anywhere in the platform.

REQUIRED BY THE SPECIFICATION

Operations Scan at 795 euro, Cyber Scan at 1,595 euro and Elite Scan at 3,295 euro, all excluding VAT, with paid add-ons. COM-001 states plainly: there is no free AI Scan product.

Scoring

IN PRODUCTION TODAY

An overallScore integer column on the scans table, and per-dimension scores constrained between 0 and 100 in the scoring engine.

REQUIRED BY THE SPECIFICATION

PRD-006 and ADD-EVAL-003 prohibit overall scores, maturity scores, readiness percentages and confidence percentages outright. Confidence becomes qualitative: High, Moderate, Low or Insufficient.

The report

IN PRODUCTION TODAY

Five sections: executive summary, dimension breakdown, detected opportunities, operational roadmap and important notes.

REQUIRED BY THE SPECIFICATION

A fixed 21-section canonical architecture rendered through a controlled design system, with a mandatory typography approval gate, internal cross-references, PDF bookmarks and a portal book viewer.

Who approves the result

IN PRODUCTION TODAY

The model generates the report and it is delivered. There is no human review step anywhere in the flow.

REQUIRED BY THE SPECIFICATION

A mandatory Expert Review workspace, an assigned reviewer, a review completion gate, a reviewer handoff package and a separate Super Admin publication action. Reviewer approval explicitly must not publish.

What has no equivalent at all

Beyond the four contradictions above, the specifications require whole capabilities that have no counterpart in the existing implementation. These are not modifications. There is nothing to modify.

- A payment and entitlement system. No payment processor is integrated anywhere in the platform today. The specification makes a verified payment the only public route to an account.
- An adaptive assessment engine with persistent structured memory, a depth engine, contradiction detection and evidence-state tracking across six states.
- An Expert Review workflow: reviewer assignment, a review brief across thirteen attention categories, a finding inspector, controlled edits with change history, targeted regeneration and a completion gate.
- A publication system with server-side binding, atomic release, immutable published versions and a delivered date derived only from successful publication.
- A translation and language pipeline producing Dutch and English review renderings with semantic equivalence validation and stale-rendering detection.
- An Executive Briefing: a premium audiovisual deliverable with professional voice-over, motion composition, captions and version binding to the approved report.
- A portal book viewer that presents the report as a digital publication rather than an embedded PDF, with responsive single and double page logic.
- An authoritative state machine of thirty named states, a logical data model of roughly thirty-five distinct concepts, an audit event catalogue and a failure and recovery matrix.

The specification anticipates exactly the argument that this is already built, and answers it in its own Definition of Done:

MASTER SPECIFICATION, DONE-001

The system is NOT Done merely because screens exist, Stripe works, AI asks questions, a PDF generates, a reviewer can click Approve or a report appears in the Client Portal.

We would not have phrased it more clearly ourselves.

What the three agreed milestones actually covered

The three milestone plan had one purpose: move the platform off the previous provider, secure it, and deploy it. It was infrastructure and migration work. It contained no product development.

Milestone by milestone

Milestone 1, tasks RM-01 to RM-63. Repository cleanup, secrets rotation, environment cleanup, GitHub migration, CI foundation, code hygiene, bug fixes, Supabase project setup, database migration from the previous provider, authentication migration, RBAC rebuild, route guards, row level security and environment separation.

Milestone 2. Storage migration to Supabase, full removal of the previous provider's dependencies, email productionisation, and verification of the core workflows that already existed.

Milestone 3, tasks RM-64 to RM-119. Notification event catalogue, deployment, security hardening, testing and quality assurance, performance and final verification.

Where the AI Scan sits in that history

This is the part that matters most for the question being asked, and it is checkable.

The AI Scan was not built during the three milestones. It already existed when Milestone 1 began, inherited from the previous build. Milestone 1 task RM-35 was titled Fix admin AI Scans KPI and describes correcting a dashboard count that read from an `ai\_scans` table which was already there. Milestone 2 migrated the AI Scan's language model call away from the previous provider's gateway to a directly configured endpoint, and migrated its report storage to Supabase.

So the three milestones migrated and repaired an AI Scan that was already in the product. They never scoped, designed or built it as a commercial assessment product, because that was not what the plan was for.

The structural test

Every task in the three agreed milestones carries a reference number, running RM-01 through RM-119. That numbering is how the plan was tracked from the beginning, in documents shared throughout.

Nothing in the new specifications has one. Not the adaptive engine, not the payment system, not Expert Review, not publication, not the Executive Briefing, not the report design system, not the telephony layer, not the IVR. The absence is not an oversight in our tracking; it is because no workstream in any of the three milestones contained them.

The scale of what has been specified

Eight specification documents have now been received across two rounds, totalling 133 pages.

Received 17 and 18 September

- AI Receptionist Functional Specification, 13 pages, 57 sections
- AI Receptionist Conversation and Response Playbook, 17 pages, 71 sections
- IVR and Telephony Functional Specification, 14 pages, 47 sections
- Scheduling Agent Functional Specification, 16 pages, 54 sections
- AI Sales Outbound Functional and Technical Specification, 21 pages, 101 sections

Received 26 September

- AI Scan Engine and Universal Assessment System, Master Specification, 31 pages, approximately 294 requirement identifiers
- AI Scan Operational, Administration and Implementation Clarification Addendum, 12 pages, approximately 127 requirement identifiers
- AI Scan Report Design System and Visual Consistency Specification, 9 pages, 23 acceptance criteria

Together these produce 302 implementation tasks. The AI Scan documents alone carry 77 separate acceptance conditions: 30 in the Master Specification, 24 in the Addendum and 23 in the Design System. Several of those are negative tests that must fail rather than pass.

For comparison, Milestone 3, the largest of the three agreed milestones, contained 56 tasks in total.

A quarter of the work is not even the named products

Of the 302 tasks, 53 are neither AI Scan nor a voice product. They are shared platform engineering that the specifications require and that no quote for four AI products would contain: a model routing layer with cost governance, a controlled learning engine, a configuration publish and rollback lifecycle, permission-aware global search across the platform, granular technical operator scopes, a rewrite of account creation across all five roles, outbound email infrastructure on a separate domain, and the acceptance evidence obligations.

One of those deserves particular mention because it reaches beyond this work entirely. The Scheduling specification requires that every supported language convey the same approved meaning through professional localisation, and states that raw machine translation is not an acceptable publication method. It applies that rule platform-wide, to the website and platform copy as well as the agents. The platform currently ships ten languages.

What we have already absorbed without raising it

A scope note that counts in only one direction is not much use, so here is the other side. The following was delivered in the last round and none of it was in the agreed plan either. We did not invoice for any of it.

- The multi-factor authentication fix that restored your access to the Admin and Super Administrator areas. That was a genuine fault on our side and we treated it as such.
- The Data Subject Rights backend, built to your specification including the requirement that authority come from a named individual grant rather than a job title.
- Profile photos across all five roles, with storage rules enforced at the database level and uploads validated by file contents rather than by what the browser claims.
- Moving the production application into the European Economic Area. Your database was in Ireland but the application was running in the United States, which meant personal data was processed outside the area before reaching a European database.
- Correcting the telephone fields on both public forms, and bringing transactional email onto your exact brand colours.
- Repairing the mobile navigation, which turned out to be a deeper fault: the site header had not been staying fixed on any device. The same fix recovered twenty eight styling rules that had never rendered in production.
- Answering the fifteen section GDPR questionnaire in full, which required auditing the schema, the provider configuration and the client side tracking.

We mention this not to present a bill for it, but because it is the honest context for the position in this note. We absorb work regularly and without complaint. What we cannot absorb is a body of work several times the size of the original engagement.

What we recommend

None of this is a refusal. Every item in these eight documents is work we want to do, and the specifications are unusually well written. Several of the harder requirements, particularly around backend authorisation, no false success, recognition not being verification, and evidence never being silently upgraded to fact, will make the system genuinely safer.

What we are asking for is that the work be scoped and sequenced as what it is, rather than absorbed into a plan that was written for a different purpose and is now complete.

Four practical suggestions

- Treat the shared foundation as its own stage. Twenty three tasks unlock all five products at once. Building them in parallel without it means building the same telephony, speech, permission and audit groundwork five separate times, which is more expensive for IO SKY, not less.
- Sequence the AI Scan ahead of the voice products. It has the shortest path to revenue at 795 to 3,295 euro per scan, it needs no telephony provider, and the payment integration it requires is also what unblocks the account creation rule that the Scheduling specification depends on.
- Scope the four voice products together rather than one at a time. They share one telephony and speech foundation. Quoting them separately means paying for that foundation more than once.
- Tell us which single capability the companies you mentioned are waiting on. That one answer changes the order we work in more than any other piece of information you could give us, and it is free to provide.

What would unblock us immediately

Two credentials, both small, both outstanding since Milestone 3, and both switching finished features back on: a replacement Supabase key, and an AI provider key. The Supabase key currently in the deployment is the one that was correctly revoked after a credential exposure and never replaced, which is why all file storage is currently unable to complete.

Sources: IO_SKY_Three_Milestone_Plan.pdf and the PHASE1_CHECKLIST.md, MILESTONE2_PROGRESS.md and MILESTONE3_CHECKLIST.md trackers maintained throughout the engagement; the eight specification documents supplied by IO SKY in September 2026. Quotations are verbatim with punctuation normalised. Claims about the current implementation were verified against the running codebase and database on 26 September 2026.